

LAB 6

<https://colab.research.google.com/drive/1qzfDPeBc09iiYNruoPimcYdwEbpL0bID#sandboxMode=true&scrollTo=VZ2YjllM4Wr7>

Q1)

```
# Compute the gradient (partial derivative) of the loss function with  
regard to each parameter  
gradient = tape.gradient(loss, parameters)
```

Q2)


A neuron is just a non-linear transformation (e.g. sigmoid/ReLU) of a linear model, implying one parameter for each input dimension, + 1 for the line intercept/constant (also called neuron "bias")

With width m , height n , previous layer's filters d and account for all such filters k in the current layer.

Number of parameters = $((m * n * d) + 1) * k$, added 1 because of the bias term for each filter.

Here :

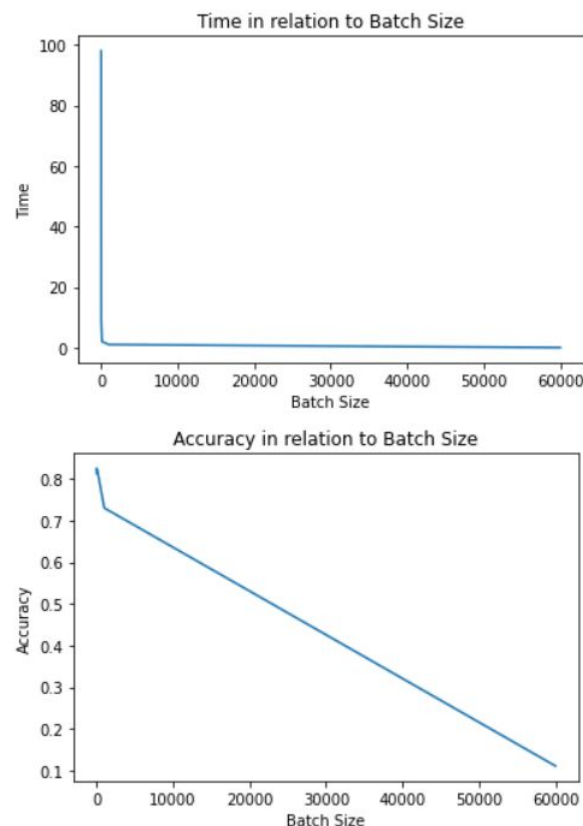
Nb parameters = $(28 * 28 + 1) * 128 = \mathbf{100480}$



Q3)

Batch size	Training times & Accuracy
1	6000/6000 [=====] 98s 2ms/step - loss: 0.5236 - accuracy: 0.8138
10	6000/6000 [=====] - 9s 2ms/step - loss: 0.4791 - accuracy: 0.8254
100	600/600 [=====] - 2s 3ms/step - loss: 0.5185 - accuracy: 0.8201
1000	60/60 [=====] - 1s 14ms/step - loss: 0.8456 - accuracy: 0.7302
60000	1/1 [=====] - 0s 2ms/step - loss: 2.3579 - accuracy: 0.1107

Plots :



We can see that with a batch size of 60000, even if the training time is almost instantaneous, the accuracy is not acceptable, so the batch size is way too large. From our observations, we can see that the **batch size of 100 seems to be the best one**. Indeed, we have the best compromise accuracy/training time, compared to the other ones.